

2020 MLB Postseason Participants Prediction

Introduction

2020 MLB season has already been hard to predict given it's short nature. This research aims to build a perfect model using team batting stats of the past eight years and then use the model to predict which teams' batting stats on 08/13/2020 is worthy of getting into postseason on a traditional 10-team postseason format.

Methods

Model was built using the combination of 16 team regular season stats: PA, R, H, 2B, 3B, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB from 2012-2019 and whether that corresponding team went into postseason or not.

Using SQL Server and Python

logistic regression
k-nearest neighbors
support vector machine
decision tree
random forest
gradient boosting
XGBoost

Step 1: Import data

import regular season stats from MLB teams who got into postseason during 2012-2019

items include Tm, PA, R, H, 2B, 3B, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB

total rows are 8(years)*10(teams each year)=80

In [1]:

```
import pandas as pd
import pyodbc

sql_conn = pyodbc.connect('''DRIVER={ODBC Driver 13 for SQL Server};
                           SERVER=ALLENHO\MSSQLSERVER002;
                           DATABASE=Playoffbound;
                           Trusted_Connection=yes''')

query = '''
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['19B$']
where Tm in ('WSN','LAD','MIL','ATL','STL','HOU','NYY','MIN','TBR','OAK')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['18B$']
where Tm in ('BOS','LAD','MIL','ATL','CHC','HOU','NYY','CLE','COL','OAK')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['17B$']
where Tm in ('BOS','LAD','COL','WSN','CHC','HOU','NYY','CLE','ARI','MIN')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['16B$']
where Tm in ('TOR','CLE','BOS','BAL','TEX','NYM','CHC','LAD','WSN','SFG')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['15B$']
where Tm in ('TOR','KCR','HOU','NYY','TEX','NYM','CHC','LAD','STL','PIT')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['14B$']
where Tm in ('BAL','KCR','OAK','LAA','DET','WSN','STL','LAD','PIT','SFG')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['13B$']
```

```

where Tm in ('BOS','TBR','OAK','CLE','DET','ATL','STL','LAD','PIT','CIN')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['12B$']
where Tm in ('TEX','BAL','OAK','NYY','DET','ATL','STL','SFG','WSN','CIN')
'''

df = pd.read_sql(query, sql_conn)

#stored as df_post
df_post = df

```

import regular season stats from MLB teams who DIDN'T get into postseason during 2012-2019

items are the same as above

total rows are 8(years)*20(teams each year)=160

In [2]:

```

sql_conn = pyodbc.connect('''DRIVER={ODBC Driver 13 for SQL Server};
                           SERVER=ALLENHO\MSSQLSERVER002;
                           DATABASE=Playoffbound;
                           Trusted_Connection=yes''')

query = '''
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['19B$']
where Tm is not null and Tm not in ('WSN','LAD','MIL','ATL','STL','HOU','NYY','MIN','TBR','OAK',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['18B$']
where Tm is not null and Tm not in ('BOS','LAD','MIL','ATL','CHC','HOU','NYY','CLE','COL','OAK',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['17B$']
where Tm is not null and Tm not in ('BOS','LAD','COL','WSN','CHC','HOU','NYY','CLE','ARI','MIN',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['16B$']
where Tm is not null and Tm not in ('TOR','CLE','BOS','BAL','TEX','NYM','CHC','LAD','WSN','SFG',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['15B$']
where Tm is not null and Tm not in ('TOR','KCR','HOU','NYY','TEX','NYM','CHC','LAD','STL','PIT',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['14B$']
where Tm is not null and Tm not in ('BAL','KCR','OAK','LAA','DET','WSN','STL','LAD','PIT','SFG',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['13B$']
where Tm is not null and Tm not in ('BOS','TBR','OAK','CLE','DET','ATL','STL','LAD','PIT','CIN',' '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['12B$']
where Tm is not null and Tm not in ('TEX','BAL','OAK','NYY','DET','ATL','STL','SFG','WSN','CIN',' '
LgAvg')
'''

df = pd.read_sql(query, sql_conn)

#stored as df_npost
df_npost = df

#add each dataframe a new column named POST, which imply whether the team made the postseason that
year
df_post['POST']= 1
df_npost['POST']= 0

#append two dataframes together
df_com=df_post.append(df_npost)

```

```
#take a look at the table we got
print(df_com)
```

	Tm	PA	R	H	HR	RBI	SB	CS	BB	SO	\
0	ATL	6302.0	855.0	1432.0	249.0	824.0	89.0	28.0	619.0	1467.0	
1	HOU	6394.0	920.0	1538.0	288.0	891.0	67.0	27.0	645.0	1166.0	
2	LAD	6282.0	886.0	1414.0	279.0	861.0	57.0	10.0	607.0	1356.0	
3	MIL	6309.0	769.0	1366.0	250.0	744.0	101.0	25.0	629.0	1563.0	
4	MIN	6392.0	939.0	1547.0	307.0	906.0	28.0	21.0	525.0	1334.0	
...	
155	PIT	6014.0	651.0	1313.0	170.0	620.0	73.0	52.0	444.0	1354.0	
156	SDP	6112.0	651.0	1339.0	121.0	610.0	155.0	46.0	539.0	1238.0	
157	SEA	6057.0	619.0	1285.0	149.0	584.0	104.0	35.0	466.0	1259.0	
158	TBR	6105.0	697.0	1293.0	175.0	665.0	134.0	44.0	571.0	1323.0	
159	TOR	6094.0	716.0	1346.0	198.0	677.0	123.0	41.0	473.0	1251.0	

	BA	OBP	SLG	OPS	TB	POST
0	0.258	0.336	0.452	0.789	2514.0	1
1	0.274	0.352	0.495	0.848	2781.0	1
2	0.257	0.338	0.472	0.810	2593.0	1
3	0.246	0.329	0.438	0.767	2429.0	1
4	0.270	0.338	0.494	0.832	2832.0	1
...
155	0.243	0.304	0.395	0.699	2138.0	0
156	0.247	0.319	0.380	0.699	2060.0	0
157	0.234	0.296	0.369	0.665	2027.0	0
158	0.240	0.317	0.394	0.711	2128.0	0
159	0.245	0.309	0.407	0.716	2231.0	0

[240 rows x 16 columns]

Take a brief look at the summary

In [3]:

```
print(df_com.describe())
```

	PA	R	H	HR	RBI	\
count	240.000000	240.000000	240.000000	240.000000	240.000000	
mean	6158.929167	712.929167	1397.520833	178.191667	679.350000	
std	97.271646	77.051138	71.511707	40.684477	75.512603	
min	5905.000000	513.000000	1199.000000	95.000000	485.000000	
25%	6085.250000	652.500000	1346.000000	148.000000	622.500000	
50%	6154.500000	709.500000	1390.500000	174.500000	675.000000	
75%	6224.250000	761.000000	1446.000000	211.000000	728.000000	
max	6475.000000	943.000000	1625.000000	307.000000	906.000000	

	SB	CS	BB	SO	BA	OBP	\
count	240.000000	240.000000	240.000000	240.000000	240.000000	240.000000	
mean	87.53750	33.195833	499.750000	1296.412500	0.252933	0.319117	
std	28.46018	8.737884	63.716652	130.141413	0.010588	0.011853	
min	19.00000	10.000000	375.000000	973.000000	0.226000	0.292000	
25%	66.00000	27.000000	452.000000	1204.000000	0.245000	0.311000	
50%	86.00000	33.000000	500.000000	1290.500000	0.252000	0.319000	
75%	105.25000	38.250000	545.500000	1384.500000	0.260000	0.327000	
max	181.00000	61.000000	656.000000	1595.000000	0.283000	0.352000	

	SLG	OPS	TB	POST
count	240.000000	240.000000	240.000000	240.000000
mean	0.409925	0.729004	2264.937500	0.333333
std	0.026815	0.036499	163.359296	0.472390
min	0.335000	0.627000	1810.000000	0.000000
25%	0.391000	0.702750	2152.500000	0.000000
50%	0.409000	0.728000	2256.500000	0.000000
75%	0.428250	0.752000	2364.000000	1.000000
max	0.495000	0.848000	2832.000000	1.000000

Take a brief look at the correlation table

In [4]:

```
df_corr = df_com.corr()
```

```
print(df_corr)
```

	PA	R	H	HR	RBI	SB	CS	\
PA	1.000000	0.728568	0.603097	0.404630	0.724551	-0.075825	-0.275590	
R	0.728568	1.000000	0.627253	0.757987	0.996604	-0.077307	-0.286587	
H	0.603097	0.627253	1.000000	0.168541	0.622822	0.010873	-0.096667	
HR	0.404630	0.757987	0.168541	1.000000	0.773527	-0.226804	-0.338698	
RBI	0.724551	0.996604	0.622822	0.773527	1.000000	-0.101979	-0.300592	
SB	-0.075825	-0.077307	0.010873	-0.226804	-0.101979	1.000000	0.563940	
CS	-0.275590	-0.286587	-0.096667	-0.338698	-0.300592	0.563940	1.000000	
BB	0.672470	0.572770	0.031033	0.490024	0.570726	-0.060325	-0.240242	
SO	0.004639	0.060778	-0.417156	0.423937	0.068943	-0.128711	-0.030335	
BA	0.512469	0.614537	0.978840	0.149331	0.609907	0.044830	-0.045988	
OBP	0.790239	0.828347	0.741023	0.438346	0.823876	-0.004032	-0.162219	
SLG	0.582243	0.922851	0.598798	0.864843	0.931042	-0.141123	-0.270421	
OPS	0.685823	0.948734	0.681851	0.779146	0.953297	-0.105815	-0.251054	
TB	0.658199	0.926030	0.685324	0.816982	0.933121	-0.148223	-0.286691	
POST	0.461904	0.488744	0.339784	0.281205	0.481147	-0.079983	-0.224696	

	BB	SO	BA	OBP	SLG	OPS	TB	\
PA	0.672470	0.004639	0.512469	0.790239	0.582243	0.685823	0.658199	
R	0.572770	0.060778	0.614537	0.828347	0.922851	0.948734	0.926030	
H	0.031033	-0.417156	0.978840	0.741023	0.598798	0.681851	0.685324	
HR	0.490024	0.423937	0.149331	0.438346	0.864843	0.779146	0.816982	
RBI	0.570726	0.068943	0.609907	0.823876	0.931042	0.953297	0.933121	
SB	-0.060325	-0.128711	0.044830	-0.004032	-0.141123	-0.105815	-0.148223	
CS	-0.240242	-0.030335	-0.045988	-0.162219	-0.270421	-0.251054	-0.286691	
BB	1.000000	0.236172	0.035724	0.647519	0.440775	0.534829	0.400129	
SO	0.236172	1.000000	-0.463401	-0.188190	0.143080	0.044561	0.101463	
BA	0.035724	-0.463401	1.000000	0.769802	0.595888	0.688866	0.655888	
OBP	0.647519	-0.188190	0.769802	1.000000	0.735855	0.866479	0.746457	
SLG	0.440775	0.143080	0.595888	0.735855	1.000000	0.975517	0.985919	
OPS	0.534829	0.044561	0.688866	0.866479	0.975517	1.000000	0.968787	
TB	0.400129	0.101463	0.655888	0.746457	0.985919	0.968787	1.000000	
POST	0.431907	-0.105900	0.356648	0.546741	0.405289	0.477010	0.403341	

	POST
PA	0.461904
R	0.488744
H	0.339784
HR	0.281205
RBI	0.481147
SB	-0.079983
CS	-0.224696
BB	0.431907
SO	-0.105900
BA	0.356648
OBP	0.546741
SLG	0.405289
OPS	0.477010
TB	0.403341
POST	1.000000

Step 2: Train the Models

There's one other consideration worth making -- the distribution of outcomes is somewhat imbalanced. Teams getting into postseason each year are less than those not. I tried an oversampling technique to see how it affected the models. Oversampling techniques are usually applied to datasets where an outcome is significantly less common. That might be a little bit of a stretch for this scenario, but I think it's worth at least checking if an oversampling technique would help. I tried fitting my different models twice -- with and without oversampling. For oversampling, I used SMOTE (synthetic minority oversampling technique).

use 4 metrics to evaluate the models, which together should give a good picture of the best overall model:

F1 score (weighted by instances of each label) ROC AUC (computed by label and weighted by frequency) balanced accuracy (for imbalanced datasets) log loss

In [5]:

```
from sklearn.preprocessing import StandardScaler

# split data into X and y
X = df_com.loc[:, 'PA': 'TB']
Y = df_com.loc[:, 'POST']
```

```

# scale and center numeric columns
X = StandardScaler().fit_transform(X)

```

In [6]:

```

from imblearn.pipeline import Pipeline
from sklearn.model_selection import cross_validate
from sklearn.metrics import f1_score, accuracy_score, log_loss, roc_auc_score, make_scorer
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier

# scoring metrics
scoring = {
    'f1_weighted': 'f1_weighted',
    'accuracy': 'balanced_accuracy',
    'roc_auc': 'roc_auc_ovr_weighted',
    'neg_log_loss': 'neg_log_loss'
}

# for results df
eval_cols = [
    'models',
    'F1 Score',
    'Balanced Accuracy',
    'ROC AUC',
    'Neg Log Loss'
]

# define classifier models
classifiers = [
    LogisticRegression(multi_class='multinomial', max_iter=10000),
    KNeighborsClassifier(),
    SVC(probability=True),
    DecisionTreeClassifier(),
    RandomForestClassifier(),
    GradientBoostingClassifier(),
    XGBClassifier()
]

# classifier names
clf_names = [
    'Logistic Regression',
    'KNN',
    'SVM',
    'Decision Tree',
    'Random Forest',
    'Gradient Boosting',
    'XGBClassifier'
]

```

C:\Users\allen\anaconda3\lib\site-packages\sklearn\externals\six.py:31: FutureWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (<https://pypi.org/project/six/>).

"(<https://pypi.org/project/six/>).", FutureWarning)

C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:144: FutureWarning: The sklearn.neighbors.base module is deprecated in version 0.22 and will be removed in version 0.24. The corresponding classes / functions should instead be imported from sklearn.neighbors. Anything that cannot be imported from sklearn.neighbors is now part of the private API.

warnings.warn(message, FutureWarning)

In [7]:

```

from imblearn.over_sampling import SMOTE
import time as time
import numpy as np

f1, acc, roc_auc, log_loss = [], [], [], []
for clf, clf_nm in zip(classifiers, clf_names):

```

```

# setup pipeline to oversample, then fit model
pipe = Pipeline([
    ('smote', SMOTE()),
    ('classify', clf)
])

start = time.time()

# cross-validate 5 times
res_smote = cross_validate(pipe, X, Y, cv=5, scoring=scoring)
results_smote = pd.DataFrame(res_smote)

stop = time.time()

print('Time to cross-validate %s = %0.3f min.' % (clf_nm, (stop - start) / 60))

# save average scores
f1.append(np.mean(results_smote.test_f1_weighted))
acc.append(np.mean(results_smote.test_accuracy))
roc_auc.append(np.mean(results_smote.test_roc_auc))
log_loss.append(np.mean(results_smote.test_neg_log_loss))

# save results to df
model_eval_smote = pd.DataFrame(data=zip(clf_names, f1, acc, roc_auc, log_loss),
                                columns=eval_cols)

display(model_eval_smote)

```

```

C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)

```

Time to cross-validate Logistic Regression = 0.009 min.

```

C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
  warnings.warn(msg, category=FutureWarning)

```

Time to cross-validate KNN = 0.003 min.


```

removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)

```

Time to cross-validate Gradient Boosting = 0.029 min.

```

C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)
C:\Users\allen\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning:
Function safe_indexing is deprecated; safe_indexing is deprecated in version 0.22 and will be
removed in version 0.24.
warnings.warn(msg, category=FutureWarning)

```

Time to cross-validate XGBClassifier = 0.015 min.

	models	F1 Score	Balanced Accuracy	ROC AUC	Neg Log Loss
0	Logistic Regression	0.727067	0.728125	0.826953	-0.533989
1	KNN	0.653692	0.659375	0.724414	-3.600277
2	SVM	0.624776	0.618750	0.739844	-0.619171
3	Decision Tree	0.651017	0.621875	0.621875	-11.800892
4	Random Forest	0.730019	0.706250	0.798828	-0.539176
5	Gradient Boosting	0.706371	0.687500	0.773828	-0.672079
6	XGBClassifier	0.712687	0.709375	0.775781	-0.654405

In [8]:

```

import numpy as np
import time as time
f1, acc, roc_auc, log_loss = [], [], [], []
for clf, clf_nm in zip(classifiers, clf_names):

    start = time.time()

    # cross-validate 5 times
    res = cross_validate(clf, X, Y, cv=5, scoring=scoring)
    results = pd.DataFrame(res)

    stop = time.time()

    print('Time to cross-validate %s = %0.3f min.' % (clf_nm, (stop - start) / 60))

    # save average scores
    f1.append(np.mean(results.test_f1_weighted))

```



```

# append(np.mean(results.test_neg_log_loss),
acc.append(np.mean(results.test_accuracy))
roc_auc.append(np.mean(results.test_roc_auc))
log_loss.append(np.mean(results.test_neg_log_loss))

# save results to df
model_eval = pd.DataFrame(data=zip(clf_names, f1, acc, roc_auc, log_loss),
                           columns=eval_cols)

display(model_eval)

```

Time to cross-validate Logistic Regression = 0.002 min.
Time to cross-validate KNN = 0.001 min.
Time to cross-validate SVM = 0.002 min.
Time to cross-validate Decision Tree = 0.001 min.
Time to cross-validate Random Forest = 0.021 min.
Time to cross-validate Gradient Boosting = 0.015 min.
Time to cross-validate XGBClassifier = 0.005 min.

	models	F1 Score	Balanced Accuracy	ROC AUC	Neg Log Loss
0	Logistic Regression	0.748284	0.715625	0.828516	-0.491554
1	KNN	0.668809	0.631250	0.739844	-2.568752
2	SVM	0.731454	0.684375	0.751563	-0.546522
3	Decision Tree	0.679634	0.650000	0.650000	-10.793494
4	Random Forest	0.713606	0.681250	0.808398	-0.537619
5	Gradient Boosting	0.735821	0.706250	0.791016	-0.624170
6	XGBClassifier	0.734542	0.703125	0.785547	-0.602285

It looks like not performing oversampling is the way to go here. The non-SMOTE'd data built models that slightly outperformed the SMOTE'd models. sklearn defines balanced accuracy as the average of recall on each class. Recall only considers false negatives and true positives -- and since SMOTE creates more data to help a model recognize minority classes, it should reduce false negatives. So, it's no surprise the SMOTE'd data performs better in balanced accuracy.

Overall, the Logistic Regression model was the best.

Step 3: Make Predictions with Logistic Regression model

import 2020 team stats as of 08/14/2020 normalized to 162 games, try to see which teams' stats on 08/13/2020 is worthy of getting into postseason on a traditional 10-team postseason format.

In [9]:

```

model = LogisticRegression(multi_class='multinomial')
model.fit(X, Y)

```

Out[9]:

```

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='multinomial', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)

```

In [10]:

```

model.coef_

```

Out[10]:

```

array([[ 0.20563564,  0.35558267, -0.43472342, -0.07174802, -0.16508224,
        -0.01143693, -0.18435707,  0.00724    , -0.18083768,  0.20025544,
         0.54154686,  0.04143439,  0.35595693, -0.26669742]])

```

In [11]:

```
df_2020=pd.read_excel(r'C:\Users\allen\Desktop\Baseball research\Postseason or bust\2020
projection for 0813.xlsx')
DF_2020=df_2020.loc[:,['PA162', 'R162', 'H162', 'HR162', 'RBI162', 'SB162', 'CS162', 'BB162', 'SO16
2', 'BA', 'OBP', 'SLG', 'OPS', 'TB162']]
DF_2020['PA']=DF_2020['PA162']
DF_2020['R']=DF_2020['R162']
DF_2020['H']=DF_2020['H162']
DF_2020['HR']=DF_2020['HR162']
DF_2020['RBI']=DF_2020['RBI162']
DF_2020['SB']=DF_2020['SB162']
DF_2020['CS']=DF_2020['CS162']
DF_2020['BB']=DF_2020['BB162']
DF_2020['SO']=DF_2020['SO162']
DF_2020['TB']=DF_2020['TB162']
DF_2020=DF_2020.loc[:, ['PA', 'R', 'H', 'HR', 'RBI', 'SB', 'CS', 'BB', 'SO', 'BA', 'OBP', 'SLG', 'OPS', 'TB']]
print(DF_2020.head())
```

	PA	R	H	HR	RBI	SB	\
0	6096.315789	750.315789	1347.157895	127.894737	707.684211	34.105263	
1	5977.800000	842.400000	1312.200000	226.800000	826.200000	64.800000	
2	6111.000000	864.000000	1467.000000	243.000000	846.000000	63.000000	
3	6096.315789	724.736842	1415.368421	196.105263	682.105263	34.105263	
4	6176.250000	840.375000	1296.000000	232.875000	789.750000	50.625000	

	CS	BB	SO	BA	OBP	SLG	OPS	TB
0	25.578947	477.473684	1219.263158	0.245	0.315	0.382	0.697	2097.473684
1	24.300000	526.500000	1644.300000	0.244	0.316	0.437	0.753	2349.000000
2	45.000000	513.000000	1305.000000	0.269	0.337	0.473	0.810	2583.000000
3	25.578947	426.315789	1492.105263	0.254	0.316	0.422	0.738	2353.263158
4	10.125000	658.125000	1620.000000	0.244	0.344	0.431	0.776	2288.250000

Try SVM model and perform hyperparameter tuning to see if it gives us better results.

In [12]:

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import GridSearchCV

model_svc = SVC(probability=True)

X_train, X_test, y_train, y_test = train_test_split(
    X, Y,
    test_size = 0.30, random_state = 101)

model_svc.fit(X_train, y_train)

preda=model_svc.predict(X_test)
print(classification_report(y_test, preda))

#-----
param_grid = {'C': [0.1, 1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
              'kernel': ['rbf']}

grid = GridSearchCV(SVC(probability=True), param_grid, refit = True, verbose = 3)

grid.fit(X_train, y_train)

grid_predictions = grid.predict(X_test)

print(classification_report(y_test, grid_predictions))
print(grid.best_params_)
print(grid.best_estimator_)
```

	precision	recall	f1-score	support
0	0.75	0.87	0.80	47
1	0.65	0.44	0.52	25
accuracy			0.72	72

macro avg	0.70	0.66	0.66	72
weighted avg	0.71	0.72	0.71	72

Fitting 5 folds for each of 25 candidates, totalling 125 fits

```
[CV] C=0.1, gamma=1, kernel=rbf .....
[CV] ..... C=0.1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=1, kernel=rbf .....
[CV] ..... C=0.1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=1, kernel=rbf .....
[CV] ..... C=0.1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=1, kernel=rbf .....
[CV] ..... C=0.1, gamma=1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=1, kernel=rbf .....
[CV] ..... C=0.1, gamma=1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.1, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.01, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.01, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.001, kernel=rbf, score=0.676, total= 0.0s
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 0.0s remaining: 0.0s
```

```
[CV] C=0.1, gamma=0.001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.001, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.001, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.0001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.0001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.0001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.0001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.0001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.0001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=0.1, gamma=0.0001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.0001, kernel=rbf, score=0.667, total= 0.0s
[CV] C=0.1, gamma=0.0001, kernel=rbf .....
[CV] ..... C=0.1, gamma=0.0001, kernel=rbf, score=0.667, total= 0.0s
[CV] C=1, gamma=1, kernel=rbf .....
[CV] ..... C=1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=1, gamma=1, kernel=rbf .....
[CV] ..... C=1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=1, gamma=1, kernel=rbf .....
[CV] ..... C=1, gamma=1, kernel=rbf, score=0.676, total= 0.0s
[CV] C=1, gamma=1, kernel=rbf .....
[CV] ..... C=1, gamma=1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=1, gamma=1, kernel=rbf .....
[CV] ..... C=1, gamma=1, kernel=rbf, score=0.667, total= 0.0s
[CV] C=1, gamma=0.1, kernel=rbf .....
[CV] ..... C=1, gamma=0.1, kernel=rbf, score=0.765, total= 0.0s
[CV] C=1, gamma=0.1, kernel=rbf .....
[CV] ..... C=1, gamma=0.1, kernel=rbf, score=0.735, total= 0.0s
[CV] C=1, gamma=0.1, kernel=rbf .....
[CV] ..... C=1, gamma=0.1, kernel=rbf, score=0.735, total= 0.0s
[CV] C=1, gamma=0.1, kernel=rbf .....
[CV] ..... C=1, gamma=0.1, kernel=rbf, score=0.735, total= 0.0s
```

[illegible]

[illegible]

```
[CV] C=1000, gamma=0.01, kernel=rbf .....
[CV] ..... C=1000, gamma=0.01, kernel=rbf, score=0.853, total= 0.0s
[CV] C=1000, gamma=0.01, kernel=rbf .....
[CV] ..... C=1000, gamma=0.01, kernel=rbf, score=0.794, total= 0.0s
[CV] C=1000, gamma=0.01, kernel=rbf .....
[CV] ..... C=1000, gamma=0.01, kernel=rbf, score=0.706, total= 0.0s
[CV] C=1000, gamma=0.01, kernel=rbf .....
[CV] ..... C=1000, gamma=0.01, kernel=rbf, score=0.788, total= 0.0s
[CV] C=1000, gamma=0.01, kernel=rbf .....
[CV] ..... C=1000, gamma=0.01, kernel=rbf, score=0.545, total= 0.0s
[CV] C=1000, gamma=0.001, kernel=rbf .....
[CV] ..... C=1000, gamma=0.001, kernel=rbf, score=0.794, total= 0.0s
[CV] C=1000, gamma=0.001, kernel=rbf .....
[CV] ..... C=1000, gamma=0.001, kernel=rbf, score=0.824, total= 0.0s
[CV] C=1000, gamma=0.001, kernel=rbf .....
[CV] ..... C=1000, gamma=0.001, kernel=rbf, score=0.676, total= 0.0s
[CV] C=1000, gamma=0.001, kernel=rbf .....
[CV] ..... C=1000, gamma=0.001, kernel=rbf, score=0.879, total= 0.0s
[CV] C=1000, gamma=0.001, kernel=rbf .....
[CV] ..... C=1000, gamma=0.001, kernel=rbf, score=0.606, total= 0.0s
[CV] C=1000, gamma=0.0001, kernel=rbf .....
[CV] .... C=1000, gamma=0.0001, kernel=rbf, score=0.853, total= 0.0s
[CV] C=1000, gamma=0.0001, kernel=rbf .....
[CV] .... C=1000, gamma=0.0001, kernel=rbf, score=0.882, total= 0.0s
[CV] C=1000, gamma=0.0001, kernel=rbf .....
[CV] .... C=1000, gamma=0.0001, kernel=rbf, score=0.706, total= 0.0s
[CV] C=1000, gamma=0.0001, kernel=rbf .....
[CV] .... C=1000, gamma=0.0001, kernel=rbf, score=0.788, total= 0.0s
[CV] C=1000, gamma=0.0001, kernel=rbf .....
[CV] .... C=1000, gamma=0.0001, kernel=rbf, score=0.636, total= 0.0s
precision recall f1-score support
```

```
0      0.73      0.87      0.80      47
1      0.62      0.40      0.49      25
```

```
accuracy          0.71      72
macro avg         0.68      0.64      0.64      72
weighted avg      0.69      0.71      0.69      72
```

```
{'C': 10, 'gamma': 0.001, 'kernel': 'rbf'}
SVC(C=10, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf',
    max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001,
    verbose=False)
```

```
[Parallel(n_jobs=1)]: Done 125 out of 125 | elapsed: 1.3s finished
```

Apparently hyperparameter tuning does not help in this case. Thus perform prediction using the old logistic regression model.

In [13]:

```
DF_2020_1 = StandardScaler().fit_transform(DF_2020)
predictions_2020_proba = model.predict_proba(DF_2020_1)
predictions_2020_pred = model.predict(DF_2020_1)
data_result = {'Team': df_2020['Tm'],
               'Probability': predictions_2020_proba[:,1]*100,
               'Prediction': predictions_2020_pred}
prediction_table = pd.DataFrame(data_result)
print(prediction_table)
```

```
Team Probability Prediction
0   ARI      31.150825         0
1   ATL      17.553153         0
2   BAL      53.163528         1
3   BOS      12.661454         0
4   CHC      83.040529         1
5   CHW      17.058149         0
6   CIN      43.136621         0
7   CLE      11.326538         0
8   COL      70.816525         1
9   DET      13.061752         0
10  HOU      60.793937         1
11  KCR       3.868289         0
12  LAA      50.827103         1
13  LAD      57.108112         1
```

13	LAN	37.100112	1
14	MIA	27.250176	0
15	MIL	4.160920	0
16	MIN	30.270392	0
17	NYM	77.672358	1
18	NYN	87.175147	1
19	OAK	47.563303	0
20	PHI	90.666869	1
21	PIT	1.140761	0
22	SDP	24.507434	0
23	SEA	10.759323	0
24	SFG	14.734669	0
25	STL	0.988297	0
26	TBR	76.495594	1
27	TEX	8.408854	0
28	TOR	1.349251	0
29	WSN	5.305377	0

In order get more accurate result. I decided to adjust my model to only include the variables that are more significantly correlated to predicting the postseason birth.

The Logistic Regression model from sklearn doesn't provide p-value automatically, so I turned to the logistic regression model from statsmodel to see which variables are less significantly correlated to predicting the postseason birth.

In [14]:

```
import statsmodels.api as sm
log_reg = sm.Logit(Y, X).fit()
print(log_reg.summary())
```

Optimization terminated successfully.

Current function value: 0.501195

Iterations 6

Logit Regression Results

```
=====
Dep. Variable:          POST    No. Observations:          240
Model:                Logit    Df Residuals:              226
Method:                MLE     Df Model:                  13
Date:                  Wed, 16 Dec 2020    Pseudo R-squ.:            0.2126
Time:                  22:21:42    Log-Likelihood:           -120.29
converged:              True     LL-Null:                  -152.76
Covariance Type:        nonrobust    LLR p-value:              6.723e-09
=====
```

	coef	std err	z	P> z	[0.025	0.975]
x1	1.3429	0.974	1.378	0.168	-0.567	3.252
x2	3.2892	2.137	1.539	0.124	-0.900	7.478
x3	-7.7348	5.924	-1.306	0.192	-19.346	3.877
x4	0.0444	0.823	0.054	0.957	-1.568	1.657
x5	-2.9628	2.257	-1.313	0.189	-7.386	1.461
x6	-0.0494	0.209	-0.237	0.813	-0.459	0.360
x7	-0.3901	0.230	-1.694	0.090	-0.842	0.061
x8	-0.3932	0.764	-0.515	0.607	-1.890	1.104
x9	-0.2961	0.239	-1.241	0.215	-0.764	0.172
x10	6.4508	4.922	1.311	0.190	-3.195	16.097
x11	-5.2292	4.448	-1.176	0.240	-13.947	3.489
x12	-17.1446	10.577	-1.621	0.105	-37.875	3.586
x13	17.7719	12.821	1.386	0.166	-7.356	42.900
x14	4.6839	8.959	0.523	0.601	-12.875	22.243

Judging from the p-value, I decided to use only those p-value under 0.18, which is PA, R, CS, SLG, OPS. A surprising discovery here is that HR(p=0.957) is wildly non-significantly correlated with postseason birth, which is a bit the contrary of what teams pursue recently.

In [15]:

```
X2 = df_com[['PA', 'R', 'CS', 'SLG', 'OPS']]
Y2 = df_com['POST']
X2 = StandardScaler().fit_transform(X2)

model2 = LogisticRegression(multi_class='multinomial')
```

```
model2.fit(X2, Y2)
```

```
Out[15]:
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='multinomial', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

```
In [16]:
```

```
DF_2020_2 = DF_2020[['PA', 'R', 'CS', 'SLG', 'OPS']]
DF_2020_2 = StandardScaler().fit_transform(DF_2020_2)
predictions_2020_proba2 = model2.predict_proba(DF_2020_2)
predictions_2020_pred2 = model2.predict(DF_2020_2)
data_result2 = {'Team': df_2020['Tm'],
                'Probability': predictions_2020_proba2[:,1]*100,
                'Prediction': predictions_2020_pred2}
prediction_table2 = pd.DataFrame(data_result2)
print(prediction_table2)
```

	Team	Probability	Prediction
0	ARI	30.528836	0
1	ATL	34.945587	0
2	BAL	50.611437	1
3	BOS	24.343055	0
4	CHC	79.979493	1
5	CHW	39.400260	0
6	CIN	22.450180	0
7	CLE	6.953635	0
8	COL	81.358858	1
9	DET	26.703734	0
10	HOU	67.514209	1
11	KCR	7.929352	0
12	LAA	51.155164	1
13	LAD	56.421238	1
14	MIA	31.429079	0
15	MIL	9.461973	0
16	MIN	38.962639	0
17	NYM	71.531353	1
18	NYY	63.583383	1
19	OAK	46.468942	0
20	PHI	64.933070	1
21	PIT	3.358363	0
22	SDP	22.172423	0
23	SEA	14.899682	0
24	SFG	18.638615	0
25	STL	1.257867	0
26	TBR	71.328421	1
27	TEX	5.664947	0
28	TOR	3.660027	0
29	WSN	2.925560	0

The list of teams were the same but there were slight differences for the probability value.

Conclusion

The result was quite satisfying given the list of teams consumes most of the powerhouse of MLB, but we still have to keep in mind that this research only considered the batting part of stats and the stats on 8/13. In my last prediction model it only considered 'PA', 'R', 'CS', 'SLG', 'OPS', which may give us a look at what matters most in teams probability of getting into postseason.

Prediction using classifiers other than Logistic Regression can be found here: https://github.com/Allen-Ho-0302/2020PostseasonPrediction-DeepLearning_XGBoost_ClassificationTree_LogisticRegression.