

Introduction

2020 MLB season has already been hard to predict given it's short nature. This research aims to build a perfect model using team batting stats of the past eight years and then use the model to predict which teams' stats on 08/13/2020 is worthy of getting into postseason on a traditional 10-team postseason format.

Methods

Model was built using the combination of 16 team regular season stats: PA, R, H, 2B, 3B, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB from 2012-2019 and whether that corresponding team went into postseason or not.

Using SQL Server and Python

Logistic Regression model

Step 1: Import data

import regular season stats from MLB teams who got into postseason during 2012-2019

items include Tm, PA, R, H, 2B, 3B, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB

total rows are 8(years)*10(teams each year)=80

In [1]:

```
import pandas as pd
import pyodbc

sql_conn = pyodbc.connect('''DRIVER={ODBC Driver 13 for SQL Server};
                           SERVER=ALLENHO\MSSQLSERVER002;
                           DATABASE=Playoffbound;
                           Trusted_Connection=yes''')

query = '''
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['19B$']
where Tm in ('WSN','LAD','MIL','ATL','STL','HOU','NYY','MIN','TBR','OAK')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['18B$']
where Tm in ('BOS','LAD','MIL','ATL','CHC','HOU','NYY','CLE','COL','OAK')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['17B$']
where Tm in ('BOS','LAD','COL','WSN','CHC','HOU','NYY','CLE','ARI','MIN')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['16B$']
where Tm in ('TOR','CLE','BOS','BAL','TEX','NYM','CHC','LAD','WSN','SFG')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['15B$']
where Tm in ('TOR','KCR','HOU','NYY','TEX','NYM','CHC','LAD','STL','PIT')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['14B$']
where Tm in ('BAL','KCR','OAK','LAA','DET','WSN','STL','LAD','PIT','SFG')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['13B$']
where Tm in ('BOS','TBR','OAK','CLE','DET','ATL','STL','LAD','PIT','CIN')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['12B$']
where Tm in ('TEX','BAL','OAK','NYY','DET','ATL','STL','SFG','WSN','CIN')
'''

df = pd.read_sql(query, sql_conn)

#stored as df_post
df_post = df
```

import regular season stats from MLB teams who DIDN'T get into postseason during 2012-2019 items are the same as above total rows are 8(years)*20(teams each year)=160

In [2]:

```
sql_conn = pyodbc.connect('DRIVER={ODBC Driver 13 for SQL Server};
                           SERVER=ALLENHO\MSSQLSERVER002;
                           DATABASE=Playoffbound;
                           Trusted_Connection=yes')

query = '''
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['19B$']
where Tm is not null and Tm not in ('WSN','LAD','MIL','ATL','STL','HOU','NYY','MIN','TBR','OAK', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['18B$']
where Tm is not null and Tm not in ('BOS','LAD','MIL','ATL','CHC','HOU','NYY','CLE','COL','OAK', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['17B$']
where Tm is not null and Tm not in ('BOS','LAD','COL','WSN','CHC','HOU','NYY','CLE','ARI','MIN', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['16B$']
where Tm is not null and Tm not in ('TOR','CLE','BOS','BAL','TEX','NYM','CHC','LAD','WSN','SFG', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['15B$']
where Tm is not null and Tm not in ('TOR','KCR','HOU','NYY','TEX','NYM','CHC','LAD','STL','PIT', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['14B$']
where Tm is not null and Tm not in ('BAL','KCR','OAK','LAA','DET','WSN','STL','LAD','PIT','SFG', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['13B$']
where Tm is not null and Tm not in ('BOS','TBR','OAK','CLE','DET','ATL','STL','LAD','PIT','CIN', '
LgAvg')
UNION ALL
select Tm, PA, R, H, HR, RBI, SB, CS, BB, SO, BA, OBP, SLG, OPS, TB
from [dbo].['12B$']
where Tm is not null and Tm not in ('TEX','BAL','OAK','NYY','DET','ATL','STL','SFG','WSN','CIN', '
LgAvg')
'''

df = pd.read_sql(query, sql_conn)

#stored as df_npost
df_npost = df

#add each dataframe a new column named POST, which imply whether the team made the postseason that
year
df_post['POST']= 1
df_npost['POST']= 0

#append two dataframes together
df_com=df_post.append(df_npost)

#take a look at the table we got
print(df_com)
```

	Tm	PA	R	H	HR	RBI	SB	CS	BB	SO	\
0	ATL	6302.0	855.0	1432.0	249.0	824.0	89.0	28.0	619.0	1467.0	
1	HOU	6394.0	920.0	1538.0	288.0	891.0	67.0	27.0	645.0	1166.0	
2	LAD	6282.0	886.0	1414.0	279.0	861.0	57.0	10.0	607.0	1356.0	
3	MIL	6309.0	769.0	1366.0	250.0	744.0	101.0	25.0	629.0	1563.0	
4	MIN	6392.0	939.0	1547.0	307.0	906.0	28.0	21.0	525.0	1334.0	
..	
155	DET	6014.0	651.0	1212.0	170.0	600.0	72.0	50.0	444.0	1254.0	

```

155 PIT 6014.0 651.0 1313.0 170.0 620.0 73.0 52.0 444.0 1354.0
156 SDP 6112.0 651.0 1339.0 121.0 610.0 155.0 46.0 539.0 1238.0
157 SEA 6057.0 619.0 1285.0 149.0 584.0 104.0 35.0 466.0 1259.0
158 TBR 6105.0 697.0 1293.0 175.0 665.0 134.0 44.0 571.0 1323.0
159 TOR 6094.0 716.0 1346.0 198.0 677.0 123.0 41.0 473.0 1251.0

```

```

      BA      OBP      SLG      OPS      TB      POST
0    0.258  0.336  0.452  0.789  2514.0      1
1    0.274  0.352  0.495  0.848  2781.0      1
2    0.257  0.338  0.472  0.810  2593.0      1
3    0.246  0.329  0.438  0.767  2429.0      1
4    0.270  0.338  0.494  0.832  2832.0      1
..     ...     ...     ...     ...     ...     ...
155  0.243  0.304  0.395  0.699  2138.0      0
156  0.247  0.319  0.380  0.699  2060.0      0
157  0.234  0.296  0.369  0.665  2027.0      0
158  0.240  0.317  0.394  0.711  2128.0      0
159  0.245  0.309  0.407  0.716  2231.0      0

```

[240 rows x 16 columns]

Step 2: Perform k-fold cross-validation on Logistic Regression Model

The expectation of repeated k-fold cross-validation is that the repeated mean would be a more reliable estimate of model performance than the result of a single k-fold cross-validation procedure. This may mean less statistical noise. One way this could be measured is by comparing the distributions of mean performance scores under differing numbers of repeats.

In [3]:

```

# compare the number of repeats for repeated k-fold cross-validation
from scipy.stats import sem
from numpy import mean
from numpy import std
from sklearn.datasets import make_classification
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
from matplotlib import pyplot

# split data into X and y
X = df_com.loc[:, 'PA':'TB']
Y = df_com.loc[:, 'POST']

# evaluate a model with a given number of repeats
def evaluate_model(X, Y, repeats):
    # prepare the cross-validation procedure
    cv = RepeatedKFold(n_splits=10, n_repeats=repeats, random_state=1)
    # create model
    model = LogisticRegression()
    # evaluate model
    scores = cross_val_score(model, X, Y, scoring='accuracy', cv=cv, n_jobs=-1)
    return scores

# configurations to test
repeats = range(1,16)
results = list()
for r in repeats:
    # evaluate using a given number of repeats
    scores = evaluate_model(X, Y, r)
    # summarize
    print('>%d mean=%.4f se=%.3f' % (r, mean(scores), sem(scores)))
    # store
    results.append(scores)
# plot the results
pyplot.boxplot(results, labels=[str(r) for r in repeats], showmeans=True)
pyplot.show()

```

```


>1 mean=0.7667 se=0.023
>2 mean=0.7625 se=0.015
>3 mean=0.7639 se=0.012
>4 mean=0.7615 se=0.010
>5 mean=0.7633 se=0.009
>6 mean=0.7618 se=0.008

```

```

>7 mean=0.7613 se=0.008
>8 mean=0.7589 se=0.008
>9 mean=0.7583 se=0.007
>10 mean=0.7579 se=0.007
>11 mean=0.7583 se=0.007
>12 mean=0.7583 se=0.007
>13 mean=0.7593 se=0.006
>14 mean=0.7598 se=0.006
>15 mean=0.7589 se=0.006

```



Ideally, we would like to select a number of repeats that shows both minimization of the standard error and stabilizing of the mean estimated performance compared to other numbers of repeats. 10 repeats seems like a good choice here.

In [4]:

```

# prepare the cross-validation procedure
cv = RepeatedKfold(n_splits=10, n_repeats=10, random_state=1)
# create model
model = LogisticRegression()
# evaluate model
scores = cross_val_score(model, X, Y, scoring='accuracy', cv=cv, n_jobs=-1)
# report performance
print('Accuracy: %.3f (%.3f)' % (mean(scores), std(scores)))

```

Accuracy: 0.758 (0.073)

Step 3: Make Predictions with Logistic Regression Model

import 2020 team stats as of 08/14/2020 normalized to 162 games, try to see which teams' stats on 08/13/2020 is worthy of getting into postseason on a traditional 10-team postseason format.

In [5]:

```

import pandas as pd
df_2020=pd.read_excel(r'C:\Users\allen\Desktop\Baseball research\Postseason or bust\2020
projection for 0813.xlsx')
df_2020=df_2020.loc[:,['PA162', 'R162', 'H162', 'HR162', 'RBI162', 'SB162', 'CS162', 'BB162', 'SO16
2', 'BA', 'OBP', 'SLG', 'OPS', 'TB162']]
df_2020['PA']=df_2020['PA162']
df_2020['R']=df_2020['R162']
df_2020['H']=df_2020['H162']
df_2020['HR']=df_2020['HR162']
df_2020['RBI']=df_2020['RBI162']
df_2020['SB']=df_2020['SB162']
df_2020['CS']=df_2020['CS162']
df_2020['BB']=df_2020['BB162']
df_2020['SO']=df_2020['SO162']
df_2020['TB']=df_2020['TB162']
DF_2020=df_2020.loc[:, ['PA', 'R', 'H', 'HR', 'RBI', 'SB', 'CS', 'BB', 'SO', 'BA', 'OBP', 'SLG', 'OPS', 'TB']]
print(df_2020.head())

```

	PA162	R162	H162	HR162	RBI162	SB162	\
0	6096.315789	750.315789	1347.157895	127.894737	707.684211	34.105263	
1	5977.800000	842.400000	1312.200000	226.800000	826.200000	64.800000	
2	6111.000000	864.000000	1467.000000	243.000000	846.000000	63.000000	

2	6111.000000	5511.000000	1107.000000	2107.000000	5107.000000	551.000000
3	6096.315789	724.736842	1415.368421	196.105263	682.105263	34.105263
4	6176.250000	840.375000	1296.000000	232.875000	789.750000	50.625000

	CS162	BB162	SO162	BA	...	PA	R	\
0	25.578947	477.473684	1219.263158	0.245	...	6096.315789	750.315789	
1	24.300000	526.500000	1644.300000	0.244	...	5977.800000	842.400000	
2	45.000000	513.000000	1305.000000	0.269	...	6111.000000	864.000000	
3	25.578947	426.315789	1492.105263	0.254	...	6096.315789	724.736842	
4	10.125000	658.125000	1620.000000	0.244	...	6176.250000	840.375000	

	H	HR	RBI	SB	CS	BB	\
0	1347.157895	127.894737	707.684211	34.105263	25.578947	477.473684	
1	1312.200000	226.800000	826.200000	64.800000	24.300000	526.500000	
2	1467.000000	243.000000	846.000000	63.000000	45.000000	513.000000	
3	1415.368421	196.105263	682.105263	34.105263	25.578947	426.315789	
4	1296.000000	232.875000	789.750000	50.625000	10.125000	658.125000	

	SO	TB
0	1219.263158	2097.473684
1	1644.300000	2349.000000
2	1305.000000	2583.000000
3	1492.105263	2353.263158
4	1620.000000	2288.250000

[5 rows x 24 columns]

In [6]:

```
### Calculate predictions: predictions
model.fit(X, Y)

predictions_2020 = model.predict(DF_2020)

predictions_2020 = [round(value) for value in predictions_2020]

print(predictions_2020)
```

[0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0]

C:\Users\allen\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Result

The model we built has a roughly 75.8% accuracy on training data set(data from 2011-2019). When trying to see which teams' stats on 08/13/2020(normalized to 162 games) are worthy of getting into postseason on a traditional 10-team postseason format, it shows BAL, CHC, COL, HOU, LAA, LAD, MIA, MIN, NYM, NYY, OAK, PHI, TBR.

Conclusion

Though the list of teams might not be exactly the powerhouse of MLB on 8/13, but we have to keep in mind that this research only considered the offense part of baseball. And it's definitely good to see teams like CHC, HOU, LAD, NYY, OAK, TBR making the list.